

Lecture 1 — February 9, 2003

Prof. Erik Demaine

Scribe: Shantonu Sen

1 Fixed-Universe Successor Problem

1.1 Motivation

Frequently, you need to store a dynamic set of n integers such that you can perform fast lookups to determine whether an element is part of the set (i.e. a *membership* test). This problem can be solved in constant time per operation using hashing, as we'll see in the next lecture. Two related pieces of information is the *successor* and *predecessor* of an integer (not necessarily in the set), that is, the two elements in the set that are immediately greater than and less than the integer, respectively. The successor and predecessor of an integer are particularly useful when the integer is not in the set, because they designate where that integer would “fit” if it were in the sorted set.

The classic solution to this problem is to maintain a balanced binary tree of the integers in your set. The membership test for a binary balanced tree can be performed in $O(\lg n)$ via binary search. The predecessor and successor functions can also be computed in $O(\lg n)$ time, by first attempting to check membership of the search element, and then moving up the tree and “leftwards” (as formalized later) for the predecessor, and moving up the tree and “rightwards” for successor.

One way to make the problem easier is to impose that the universe is not the set of all integers, but integers from 0 to $u - 1$. This assumption is called the *fixed-universe assumption*. Together with the operations described above, the problem is called the *fixed-universe successor problem* or *fixed-universe predecessor problem*. This problem is also referred to as *Interval Union-Split-Find* [MNA88] and as *priority queues* [vEKZ77].

1.2 Formal definition

We would like to create a data structure with the follow properties:

Goal: Maintain a dynamic subset S of size n from the universe $\mathcal{U} = \{0, 1, 2, 3, \dots, u - 1\}$ of size u

Supported operations:

- $\text{Insert}(x \in \mathcal{U} \notin S)$: add an element to S
- $\text{Delete}(x \in S)$: remove an element from S
- $\text{Successor}(x \in \mathcal{U})$: find the smallest element $\in S$ that is $> x$
- $\text{Predecessor}(x \in \mathcal{U})$: find the largest element $\in S$ that is $< x$

Desired performance: Better than $O(\lg n)$ for all operations, with the time bound possibly depending on u . Traditional balanced binary tree operate in $O(\lg n)$ without the fixed-universe assumption. In particular, can we achieve $O(\lg \lg u)$?

1.3 Known results

The standard solution to successor problem uses a balanced binary search trees, with a running time of $O(\lg n)$ per operation, using the comparison model on a pointer machine without the fixed-universe assumption. There are several other solutions, depending on the model of computation we consider.

1.3.1 Models


We will consider many different models of computation during this class. For starters, here are the three that have been studied for the fixed-universe successor problem:

Pointer-machine model. On a pointer machine, the data structure is described by a directed graph, where each node stores a constant number of labeled outgoing pointers and a constant number of integers. In other words, you have a constant branching factor. For the fixed-universe successor problem, there is a pointer to each element in the universe \mathcal{U} , and the input to an operation is one of these pointers.

Random Access Machine (RAM). In a RAM, memory is laid out as a finite array of slots. If you know the index of an entry, you can jump to its location and do a memory access in $O(1)$ time.

For example, below is a representation of a memory store that holds the numbers 1, 3, and 6 in the first 3 slots. Memory addresses can be loaded from and stored to, and the results of loads can be combined with arithmetic or logical operators.

The cost of an algorithm using a RAM is linear in the number of total instructions. That is, both memory accesses and arithmetic/logical operations cost $O(1)$ time.



Cell-probe model. The cell-probe model is just like a RAM, except that arithmetic/logical computation is “free”. In this model, an algorithm’s cost is linear in the number of memory accesses. This model is rather unrealistic, but is commonly used for lower bounds; any lower bound on the cell-probe model also applies to the RAM.

1.3.2 Results

Here are several upper and lower bounds for the successor and fixed-universe successor problem, and their models of computation.

- Balanced binary search trees
 - Comparison model using a pointer machine (no fixed-universe assumption)

- $O(\lg n)$ time per operation
- $O(n)$ space
- van Emde Boas [vEKZ77, vEB77] (this lecture)
 - Fixed-universe model on pointer machine (but we will describe the algorithm as working on a RAM)
 - $O(\lg \lg u)$ time per operation
 - $O(u)$ space
- Lower bound by Mehlhorn, Näher, and Alt [MNA88]
 - Pointer machine
 - Lower bound of $\Omega(\lg \lg u)$ time per operation
- y-fast trees [Wil83]
 - Randomized algorithm on a RAM
 - $O(\lg \lg u)$ time per operation
 - $O(n)$ space
- Lower bound by Beame and Fich [BF02]
 - Cell-probe model
 - Static case, no Insert/Delete
 - Lower bound of $\Omega\left(\min\left\{\frac{\lg \lg u}{\lg \lg \lg u}, \sqrt{\frac{\lg n}{\lg \lg n}}\right\}\right)$ time per operation, for any data structure using only $O(n^{O(1)})$ space
- Exponential search trees [AT99]
 - RAM model
 - $O\left(\min\left\{\frac{\lg \lg u}{\lg \lg \lg u} \lg \lg n, \sqrt{\frac{\lg n}{\lg \lg n}}\right\}\right)$ worst-case time per operation

2 $O(\lg \lg u)$ solution: van Emde Boas structure

Our goal for this lecture is to achieve $O(\lg \lg u)$ running time. Based on existing techniques for analyzing asymptotic running time of algorithms, we have some intuition about how we might end up with this type of running time.

One approach of traversing a data structure might involve doing binary search over $O(\lg u)$ things. Because a binary search runs in logarithmic time, this would yield performance of $O(\lg \lg u)$.

Another possibility is creating a recurrence relation whose solution is $O(\lg \lg u)$. For instance, consider the relation:

$$\begin{aligned}
T(u) &= T(\sqrt{u}) + O(1) \\
T'(\lg u) &= T'(\lg \sqrt{u}) + O(1) && \text{--- let } T'(\lg v) = T(v) \\
T'(\lg u) &= T'(\frac{1}{2} \lg u) + O(1) && \text{--- pull out the square root} \\
T'(x) &= T'(\frac{1}{2}x) + O(1) && \text{--- substitute } x = \lg u \\
T' &\text{ is } O(\lg x) && \text{--- using Master Method} \\
T' &\text{ is } O(\lg \lg u) && \text{--- substitute for } x = \lg u
\end{aligned}$$

The van Emde Boas structure will employ the second technique to achieve $O(\lg \lg u)$ running time. In a certain sense, it will also correspond to binary searching over the $\Theta(\lg u)$ levels of a complete binary tree on u leaves.

2.1 Starting point: precompute answers

One technique for solving the problem is to use an array to store the successor value for every element x in the universe \mathcal{U} . Similarly, store the predecessors to support Predecessor queries. For example, if our set contains the elements $\{1, 3, 7\}$, and we wish to store the successors for any possibly query element x , our array would look like:

1	3	3	7	7	7	7	...			
0	1	2	3	4	5	6	7			u-1

This makes Successor and Predecessor queries very fast, $O(1)$ time, using only a single load to fetch the answer. But updates are slow, $\Theta(n)$ time in the worst case, because we may need to update many slots to point to the newly inserted element.

2.2 Starting point: store a bit vector

Another approach is to use a RAM to store a bit vector of the elements present in the set. If the set has the elements $\{2, 3, 6, 8, 9\}$, then the RAM looks like:

0	0	1	1	0	0	1	0	1	1	
0	1	2	3							u-1

Now, queries are slower, $O(n)$ time, requiring a linear search for the next present element. However, updates are now fast, $O(1)$ time, because we just have to modify the slot corresponding to the element being added or removed.

2.3 Improving bit-vector search time with a binary tree

Using the bit-vector representation as a starting point, we build up a binary tree of OR relations. This will let us know at each node whether there are any present elements in a subtree.

To find a successor, move up the tree until you enter a node from the left and there is a 1 on the right branch. Then go down the right branch, staying as close to the left as possible while following 1 branches, until you find an element. The running time is therefore $O(\lg u)$, as it is for updates.

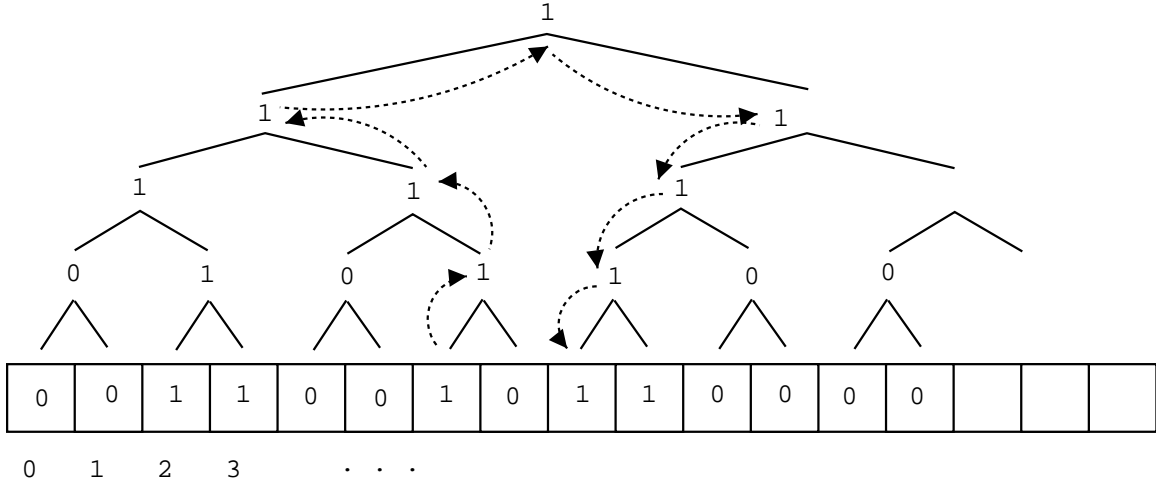


Figure 1: Bit-vector with OR tree

2.4 Variation: Use trees of constant height

Instead of a strictly binary tree, cluster the bit vector into \sqrt{u} groups of size \sqrt{u} . Let us call each cluster $\text{sub}[0], \text{sub}[1], \dots, \text{sub}[\sqrt{u}-1]$. Because each cluster has \sqrt{u} elements, $\text{sub}[i]$ represents the elements $\{i\sqrt{u}, i\sqrt{u} + 1, \dots, (i+1)\sqrt{u} - 1\} \in \mathcal{U}$.

When searching for a successor, start out in the cluster representing your query element. Do a linear search within that cluster for a successor, and if one is not present, look only at the “emptiness” summary bits for the subsequent clusters to find the next non-empty cluster. Once such a cluster is found (if there is a successor at all in the set), do a linear search within that cluster, and we are guaranteed to succeed.

Queries take $O(\sqrt{u})$ time, because we do a linear search of up to \sqrt{u} elements in two clusters, and a linear search of up to \sqrt{u} emptiness summary bits. However, inserts take $O(1)$ time, because they just need to update a bit in a cluster and possibly the emptiness summary bit of that tree.

2.5 Bit manipulation: Helper functions

Going forward, it will be useful to have some tools for manipulating the binary representation of an element $x \in \mathcal{U}$, which uses $\lceil \lg u \rceil$ bits. For example, we might express 55 in a universe of size 256 as 00110111_2 . Two functions we will want are:




Figure 2: Bit-vector grouped into trees of size \sqrt{u}

- $\text{high}(x)$ = high-order half of bits (0011₂ for the example)
- $\text{low}(x)$ = low-order half of bits (0111₂ for the example)

Expressed more mathematically, for a given value x , these functions are:

- $\text{high}(x) = \lfloor x / \sqrt{u} \rfloor$
- $\text{low}(x) = x \bmod \sqrt{u}$

If we consider our clusters above, these two functions really give us a way to map an element in the universe into a bucket, and more specifically a location within a given bucket

- $\text{high}(x)$ = which of the \sqrt{u} clusters x is in
- $\text{low}(x)$ = index within that cluster

2.6 Refined attempt: use recursion

If we can recursively apply the \sqrt{u} solution, we should be able to get down to $O(\lg \lg u)$ time per operation.

View the universe \mathcal{U} as a *structure* of size u . In general, suppose we have a set of elements represented by a structure S of size $|S|$. Split it into $\sqrt{|S|}$ substructures each of size $\sqrt{|S|}$. Each substructure (cluster) is named $\text{sub}[S][0]$, $\text{sub}[S][1]$, \dots , $\text{sub}[S][\sqrt{|S|} - 1]$.

Because we need to store which substructures are empty, we can recursively use another substructure of size $\sqrt{|S|}$ called $\text{summary}[S]$. Each element in this summary structure will correspond to the emptiness of one of the $\sqrt{|S|}$ substructures.

Insertion therefore corresponds to two recursive calls, one in the appropriate substructure, and possibly one in the summary structure:

$\text{Insert}(x, S)$:

- Insert ($\text{low}(x)$, $\text{sub}[S][\text{high}(x)]$)
- Insert ($\text{high}(x)$, $\text{summary}[S]$) if $\text{sub}[S][\text{high}(x)]$ was empty

Recurrence relation for Insert's running time:

$$\begin{aligned}
T(u) &= 2T(\sqrt{u}) + O(1) \\
T'(\lg u) &= 2T'(\lg \sqrt{u}) + O(1) \\
T'(\lg u) &= 2T'(\frac{1}{2} \lg u) + O(1) \\
T'(x) &= 2T'(\frac{1}{2}x) + O(1) \\
T' &\text{ is } O(x) \\
T' &\text{ is } O(\lg u)
\end{aligned}$$

Oops, we have too many recursive calls to get $O(\lg \lg u)$.

Successor is similar to the \sqrt{u} solution, but with recursive calls. First we look in the appropriate substructure. If we don't find the element there, we look in the summary structure for the next next nonempty substructure. The key observation is that this operation is another successor query. Then we find the minimum element in that substructure, which can be viewed as finding the successor of $-\infty$.

Successor(x, S):

```

j ← Successor(low(x), sub[S][high(x)])
if j < ∞
    return j + high(x) · √|S|
i ← Successor(high(x), summary[S])
j ← Successor(−∞, sub[S][i])
return j + i · √|S|

```

Recurrence relation for Successor:

$$\begin{aligned}
T(u) &= 3 * T(\sqrt{u}) + O(1) \\
T'(\lg u) &= 3 * T'(\lg \sqrt{u}) + O(1) \\
T'(\lg u) &= 3 * T'(\frac{1}{2} \lg u) + O(1) \\
T' &\text{ is } O(\lg u)^{\lg 3}
\end{aligned}$$

This is much worse than we want, even worse than logarithmic. The reason is that we are making too many recursive calls. We need to reduce our running time by reducing the number of recursive calls from 2 and 3 down to 1.

2.7 Store min and max of each cluster to reduce recursion

We can reduce the number of recursive calls to Insert or Successor by being mindful of when we can precalculate information without needing to do a full-blown search. One mechanism is to

cache the minimum and maximum element contained by each structure S as an additional piece of information associated with the structure S . We will refer to these as $\min[S]$ and $\max[S]$.

Accessing the minimum element in a structure S is now $O(1)$. If we can replace one or more of the recursive calls in `Insert` or `Successor` with a quick access to the minimum element of the structure, we can reduce the overall running time of the algorithms dramatically.

First, let us assess how this helps `Successor`:

```
Successor( $x, S$ ):
  if  $\text{low}(x) < \max[\text{sub}[S][\text{high}(x)]]$ 
     $j \leftarrow \text{Successor}(\text{low}(x), \text{sub}[S][\text{high}(x)])$ 
    return  $\text{high}(x) \cdot \sqrt{|S|} + j$ 
  else
     $i \leftarrow \text{Successor}(\text{high}(x), \text{summary}[S])$ 
    return  $\min[\text{sub}[S][i]] + i \cdot \sqrt{|S|}$ 
```

We know what cluster x should appear in ($\text{high}(x)$), and if there's an element larger than x in that cluster, we just need to search in that cluster. If there was no such element, we can look at the summary structure for the next non-empty cluster, and return the smallest element in that cluster.

Since the condition of the “if” can be calculated in $O(1)$ time, and either branch makes only one recursive call to `Successor` on a structure of size $\sqrt{|S|}$, we've succeeded in reducing the amount of recursion to 1. The recurrence relation now looks like:

$$\begin{aligned} T(u) &= T(\sqrt{u}) + O(1) \\ T'(\lg u) &= T'(\lg \sqrt{u}) + O(1) \\ T'(\lg u) &= T'(\frac{1}{2} \lg u) + O(1) \\ T' &\text{ is } O(\lg \lg u) \end{aligned}$$

Even though `Successor` is now performing well, this strategy does not work well for `Insert`. Fundamentally, `Insert` would still require 2 recursive calls to `Insert`—one to `Insert` into a substructure, and one to update the `summary[S]` structure. In order to have a running time of $O(\lg \lg u)$, we need to eliminate one of the recursive calls.

2.7.1 Final Solution: van Emde Boas structure

Our solution to this dilemma is to just not recurse unless absolutely needed. We do this by using $\min[S]$ as a sort of very-cheap ($O(1)$ -time) very-small (single-element) cache for the data structure. If a structure holds only a single element, that element is non-recursively stored in the $\min[S]$ slot.

Now, to check whether an entire structure S is empty, we can just check whether $\min[S]$ is unset, instead of traversing the summary structure. More importantly, inserting into an empty structure is also a $O(1)$ -time operation, because we can just set the \min value and be done.

The final forms of our `Insert` and `Successor` functions are as follows. Note that both make only a




Figure 3: Representation of full van Emde Boas structure

single recursive call to themselves, so using the derivation above for calculating running time, both these functions run in $O(\lg \lg u)$ time.

Insert(x, S):

```

if  $x < \min[S]$  then swap  $x$  &  $\min[S]$ 
if  $\text{sub}[S][\text{high}(x)]$  is empty:
    Insert ( $\text{high}(x)$ ,  $\text{summary}[S]$ )
     $\min[\text{sub}[S][\text{high}(x)]] \leftarrow \text{low}(x)$ 
else
    Insert ( $\text{low}(x)$ ,  $\text{sub}[S][\text{high}(x)]$ )
if  $x > \max[S]$  then  $\max[S] \leftarrow x$ 

```

Successor(x, S):

```

if  $\text{low}(x) < \max[\text{sub}[S][\text{high}(x)]]$ :
     $j \leftarrow \text{Successor}(\text{low}(x), \text{sub}[S][\text{high}(x)])$ 
    return  $\text{high}(x) \cdot \sqrt{|S|} + j$ 
else
     $i \leftarrow \text{Successor}(\text{high}(x), \text{summary}[S])$ 
    return  $\min[\text{sub}[S][i]] + i \cdot \sqrt{|S|}$ 

```

References

- [AT99] Arne Andersson and Mikkel Thorup. Tight(er) worst-case bounds on dynamic searching and priority queues. In *Proceedings of the 32nd Annual ACM symposium on Theory of computing*, pages 335–342. ACM, 1999.
- [BF02] Paul Beame and Faith E. Fich. Optimal bounds for the predecessor problem and related problems. *Journal of Computer and System Sciences*, 65(1):38–72, August 2002.
- [MNA88] Kurt Mehlhorn, Stefan Näher, and Helmut Alt. A lower bound on the complexity of the union-split-find problem. *SIAM Journal on Computing*, 17(6):1093–1102, December 1988.
- [vEB77] P. van Emde Boas. Preserving order in a forest in less than logarithmic time and linear space. *Information Processing Letters*, 6(3):80–82, 1977.
- [vEKZ77] Peter van Emde Boas, R. Kaas, and E. Zijlstra. Design and implementation of an efficient priority queue. *Math. Systems Theory*, 10:99–127, 1977.
- [Wil83] Dan E. Willard. Log-logarithmic worst case range queries are possible in space $\theta(n)$. *Information Processing Letters*, 17(2):81–84, 1983.

Lecture 2 — February 12, 2003

*Prof. Erik Demaine**Scribe: Jeff Lindy*

1 Overview

In the last lecture we considered the successor problem for a bounded universe of size u . We began looking at the van Emde Boas [3] data structure, which implements Insert, Delete, Successor, and Predecessor in $O(\lg \lg u)$ time per operation.

In this lecture we finish up van Emde Boas, and improve the space complexity from our original $O(u)$ to $O(n)$. We also look at perfect hashing (first static, then dynamic), using it to improve the space complexity of van Emde Boas and to implement a simpler data structure with the same running time, y -fast trees.

2 van Emde Boas

2.1 Pseudocode for vEB operations

We start with pseudocode for Insert, Delete, and Successor. (Predecessor is symmetric to Successor.)

Insert(x, S)

```
if  $x < \text{min}[S]$ : swap  $x$  and  $\text{min}[S]$ 
if  $\text{min}[\text{sub}[S][\text{high}(x)]] = \text{nil}$  : // was empty
    Insert( $\text{low}(x)$ ,  $\text{sub}[S][\text{high}(x)]$ )
     $\text{min}[\text{sub}[S][\text{high}(x)]] \leftarrow \text{low}(x)$ 
else:
    Insert( $\text{high}(x)$ ,  $\text{summary}[S]$ )
if  $x > \text{max}[S]$ :  $\text{max}[S] \leftarrow x$ 
```

Delete(x, S)

```
if  $\text{min}[S] = \text{nil}$  or  $x < \text{min}[S]$ : return
if  $\text{min}[S] = x$ :
     $i \leftarrow \text{min}[\text{summary}[S]]$ 
```

```

 $x \leftarrow i\sqrt{|S|} + \min[\text{sub}[S][i]]$ 
 $\min[S] \leftarrow x$ 
Delete(low( $x$ , sub[W][high( $x$ )))
if  $\min[\text{sub}[S][\text{high}(x)]] = \text{nil}$  : // now empty
    Delete(high( $x$ ), summary[S])
// in this case, the first recursive call was cheap

```

```

Successor( $x$ ,  $S$ )
if  $x < \min[S]$ : return  $\min[S]$ 
if low( $x$ ) < max[sub[S][high(x)]]
    return high( $x$ )  $\sqrt{|S|} + \text{Successor}(\text{low}(x), \text{sub}[S][\text{high}(x)])$ 
else:
     $i \leftarrow \text{Successor}(\text{high}(x), \text{summary}[S])$ 
    return  $i\sqrt{|S|} + \min[\text{sub}[S][i]]$ 

```

2.2 Tree view of van Emde Boas

The van Emde Boas data structure can be viewed as a tree of trees.

The upper and lower “halves” of the tree are of height $\frac{1}{2} \lg u$, that is, we are cutting our tree in halves by level. The upper tree has \sqrt{u} nodes, as does each of the subtrees hanging off its leaves. (These subtrees correspond to the *sub*[*S*] data structures in Section 2.1.)

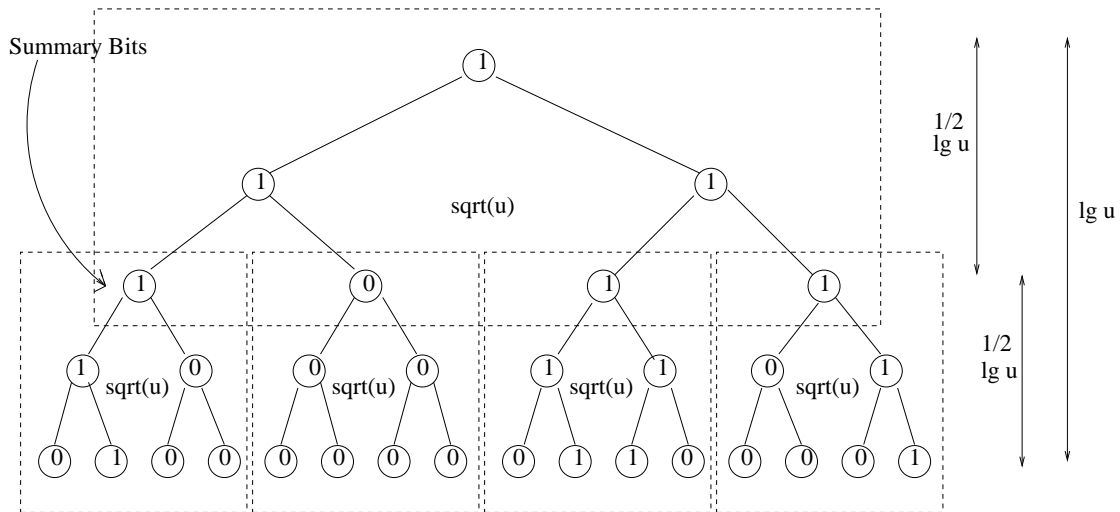


Figure 1: In this case, we have the set 1, 9, 10, 15.

We can consider our set as a bit vector (of length 16) at the leaves. Mark all ancestors of present leaves. The function $\text{high}(x)$ tells you how to walk through the upper tree. In the same way, $\text{low}(x)$ tells you how to walk through the appropriate lower tree to reach x .

For instance, consider querying for 9. $9_{10} = 1001_2$ which is equivalent to “right left left right”; $\text{high}(9) = 10_2$ (“right left”), $\text{low}(9) = 01_2$ (“left right”).

2.3 Reducing Space in vEB

It is possible that this tree is wasting a lot of space. Suppose that in one of the subtrees there was nothing whatsoever?

Instead of storing all substructures/lower trees explicitly, we can store $\text{sub}[S]$ as a dynamic hash table and store only nonempty substructures.

We perform an amortized analysis of space usage:

1. *Hash* – Charge space for hashtable to nonempty substructures, each of which stores an element as min.
2. *Summary* – Charge elements in the summary structure to the mins of corresponding substructures.

When a substructure has just one element, stored in the min field, it does not store a hashtable, so that it can occupy just $O(1)$ space. In all, the tree as described above uses $O(n)$ space, whereas van Emde Boas takes more space, $O(u)$, as originally considered [3]. For this approach to work, though, we need the following hashtable black box:

BLACKBOX for hash table:

Insert in $O(1)$ time, amortized and expected

Delete in $O(1)$ time, amortized and expected

Search in $O(1)$ time, worse case

Space is $O(n)$

3 Transdichotomous Model

There are some issues with using the comparison model (too restrictive) and the standard unit cost RAM (too permissive, in that we can fit arbitrarily large data structures into single words). Between the two we have the transdichotomous model.

This machine is a modified RAM, and is also called a *word RAM*, with some set size word upon which you can perform constant time arithmetic.

We want to bridge problem size and machine model. So we need a word that is big enough ($\lg n$) to index everything we are talking about. (This is a reasonable model since, say, a 64 bit word can index a truly enormous number of integers, 2^{64} .) With this model, we can manipulate a constant number of words in a constant time.

4 Perfect Hashing

Perfect hashing is a way to implement hash tables such that the expected number of collisions is minimized. We do this by building a hash table of hash tables, where the choice of hash functions for the second level can be redone to avoid clobbering elements through hash collision.

We'll start off with the static case; we already have all the elements, and won't be inserting new ones.

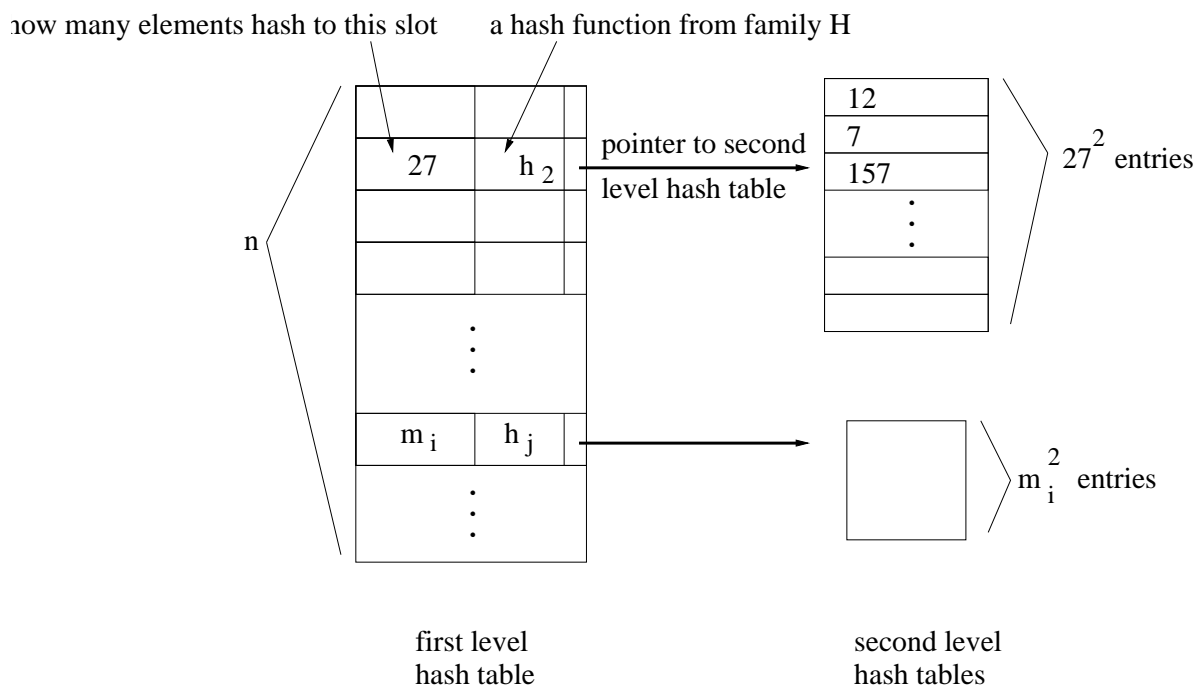


Figure 2: We're storing n elements from our universe U . h_1 maps some $u \in U$ into the first level hash table which is of size n . Each entry in the first level hash table has a count of how many elements hash there, as well as a second hash function h_i which takes elements to a spot in the appropriate second level hash table.

A set H of hash functions of the form $h : U \rightarrow \{0, 1, \dots, m - 1\}$ is called *universal* if

$$|\{h \in H : h(x) = h(y)\}| = |H|/m \quad \text{for all } x, y \in U, x \neq y$$

That is, the probability of a hash collision is equal to the size of the universe over the size of the table.

Universal sets of hash functions exist. For instance, take key k , and write it in base m :

$$k = \langle k_0, k_1, \dots, k_r \rangle$$

Choose some other vector of this size at random:

$$a = \langle a_0, a_1, \dots, a_r \rangle$$

Then the hash function

$$h_a = \sum_{i=0}^r a_i k_i \bmod m$$

is essentially weighting each k_i , via a dot product. So letting the vector a vary, we have a family of hash functions, and it turns out to be universal.

Suppose m_i elements hash to bucket i (a slot in the first hash table).

The expected number of collisions in a bucket is

$$E[\# \text{ of collisions in bucket } i] = \sum_{x \neq y} Pr\{x \text{ and } y \text{ collide}\}$$

Because the second table is of size m_i^2 , this sum is equal to

$$\binom{m_i}{2} \frac{1}{m_i^2} < \frac{1}{2} \quad (\text{by Markov})$$

The expected total size of all second-level tables is

$$E[\text{space}] = E \left[\sum_{i=0}^{n-1} m_i^2 \right]$$

Because the total number of conflicting pairs in the first level can be computed by counting pairs of elements in all second-level tables, this sum is equal to

$$n + 2 E \left[\sum_{i=0}^{n-1} \binom{m_i}{2} \right] = n + \frac{2 \binom{n}{2}}{n} < 2n$$

So our expected space is about $2n$. More specifically, the probability $Pr[\text{space} > 4n] < \frac{1}{2}$.

If we clobber an element, we pick some other hash function for that bucket. We will need to choose new hash functions only an expected constant number of times per bucket, because we have a constant probability of success each time.

4.1 Dynamic Perfect Hashing

Making a perfect hash table dynamic is relatively straightforward, and can be done by keeping track of how many elements we have inserted or deleted, and rebuilding (shrinking or growing) the entire hash table. This result is due to Dietzfelbinger et al. [4].

Suppose at time t there are n^* elements.

Construct a static perfect hash table for $2n^*$ elements. If we insert another n^* elements, we rebuild again, this time a static perfect hash table for $4n^*$ elements.

As before, whenever there's a collision at the second-level hash table, just rebuild the entire second-level table with a new choice of hash function from our family H . This won't happen very often, so amortized the cost is low.

If we delete elements until we have $n^*/4$, we rebuild for n^* instead of $2n^*$. (If we delete enough, we rebuild at half the size.)

All operations take amortized $O(1)$ time.

5 y -fast Trees

5.1 Operation Costs

y -fast trees [1, 2] accomplish the same running times as van Emde Boas— $O(\lg \lg u)$ for Insert, Delete, Successor, and Predecessor—but they are simple once you have dynamic perfect hashing in place.

We can define correspondences $u \in U \longleftrightarrow \text{bit string} \longleftrightarrow \text{path in a balanced binary search tree}$.

For all x in our set, we will store a dynamic perfect hash table of all prefixes of these strings. Unfortunately, this structure requires $O(n \lg u)$ space, because there are n elements, each of which have $\lg u$ prefixes. Also, insertions and deletions cost $\lg u$. We'll fix these high costs later.

What does this structure do for Successor, though? To find a successor, we binary search for the lowest “filled” node on the path to x . We ask, “Is half the path there?” for successively smaller halves. It takes $O(\lg \lg u)$ time to find this node.

Suppose we have each node store the min and max element for its subtree (in the hash table entry). This change adds on some constant cost for space. Also, we can maintain a linked list of the entire set of occupied leaves. Then we can find the successor and the predecessor of x in $O(1)$ time given what we have so far.

Using this structure, we can perform Successor/Predecessor in $O(\lg \lg u)$ time, and Insert/Delete in $O(\lg u)$ time, using $O(n \lg u)$ space.

5.2 Indirection

We cluster the n present elements into consecutive groups of size $\Theta(\lg u)$.

- Within a group, we use a balanced binary search tree to do Insert/Delete/Successor/Predecessor in $O(\lg \lg u)$ time.
- We use hash table as above to store one representative element per group $\Rightarrow O(n)$ space
- If a search in the hashtable narrows down to 2 groups, look in both.
- Insert and Delete are delayed by the groups, so it works out to be $O(1)$ amortized to update the hash table, and $O(\lg \lg u)$ total.




Figure 3: Using indirection in y -fast trees.

References

- [1] Dan E. Willard, “Log-logarithmic worst-case range queries are possible in space $\Theta(n)$,” *Information Processing Letters*, 17:81–84. 1984.
- [2] Dan E. Willard, “New trie data structures which support very fast search operations,” *Journal of Computer and System Sciences*, 28(3):379–394. 1984.
- [3] P. van Emde Boas, “Preserving order in a forest in less than logarithmic time and linear space,” *Information Processing Letters*, 6:80–82, 1977.
- [4] M. Dietzfelbinger, A. Karlin, K. Mehlhorn, F. Meyer auf der Heide, H. Rohnert, and R. E. Tarjan, “Dynamic perfect hashing: upper and lower bounds”, *SIAM Journal on Computing*, 23:738–761, 1994. <http://citeseer.nj.nec.com/dietzfelbinger90dynamic.html>.